

Project: Vehicle-Detection
Name: Dongdong Wu

Overview:

Vehicle detection in video-stream or in image is popular issue of Computer vision. The main idea is, by using pre-trained model, detect vehicle object. As a classification problem, the image are divided into small patches, each of which will be run through a classifier to determine whether there are objects in the patch. Then the bounding boxes will be assigned to locate around patches that are classified with high probability of present of an object. In the regression approach, the whole image will be run through a convolutional neural network to directly generate one or more bounding boxes for objects in the images.


During this project, we will use tiny-YOLO which is based off of the Darknet reference network and is much faster but less accurate than the normal YOLO model. For real time vehicle detection, speed(FPS) are critical import comparing to detection on image. By comparing model's feature:

| Model | Train | Test | mAP | FLOPS | FPS | Cfg | Weights |
|----------------|---------------|----------|------|----------|-----|-----|-------------------------|
| Old YOLO | VOC 2007+2012 | 2007 | 63.4 | 40.19 Bn | 45 | | link |
| SSD300 | VOC 2007+2012 | 2007 | 74.3 | - | 46 | | link |
| SSD500 | VOC 2007+2012 | 2007 | 76.8 | - | 19 | | link |
| YOLOv2 | VOC 2007+2012 | 2007 | 76.8 | 34.90 Bn | 67 | cfg | weights |
| YOLOv2 544x544 | VOC 2007+2012 | 2007 | 78.6 | 59.68 Bn | 40 | cfg | weights |
| Tiny YOLO | VOC 2007+2012 | 2007 | 57.1 | 6.97 Bn | 207 | cfg | weights |
| SSD300 | COCO trainval | test-dev | 41.2 | - | 46 | | link |
| SSD500 | COCO trainval | test-dev | 46.5 | - | 19 | | link |
| YOLOv2 608x608 | COCO trainval | test-dev | 48.1 | 62.94 Bn | 40 | cfg | weights |
| Tiny YOLO | COCO trainval | - | - | 7.07 Bn | 200 | cfg | weights |

That is why I choose this model to do detection.

Image Pre-process:

Since Yolo model's input format limit, I crop and resize the image to 448*448 format. Then do normalization to {-1, 1} range. The example image :



Model Architecture:

During this project, i use keras as API to implement yolo-model. According to pre-trained yolo model, this model consist of 9 convolution layers and 3 full connected layers. Each convolution layer consists of convolution, relu and max-pooling operations. The first 9 convolution layers can be understood as the feature extractor, whereas the last three full connected layers can be understood as the "regression head" that predicts the bounding boxes. The model summary reports:

The detail information can be found:

<https://pjreddie.com/darknet/yolov1/>

| Layer (type) | Output Shape | Param # | Connected to |
|---------------------------------|----------------------|---------|-----------------------|
| ===== | | | |
| convolution2d_1 (Convolution2D) | (None, 16, 448, 448) | 448 | |
| convolution2d_input_1[0][0] | | | |
| ----- | | | |
| leakyrelu_1 (LeakyReLU) | (None, 16, 448, 448) | 0 | convolution2d_1[0][0] |
| ----- | | | |
| maxpooling2d_1 (MaxPooling2D) | (None, 16, 224, 224) | 0 | leakyrelu_1[0][0] |
| ----- | | | |
| convolution2d_2 (Convolution2D) | (None, 32, 224, 224) | 4640 | maxpooling2d_1[0][0] |
| ----- | | | |
| leakyrelu_2 (LeakyReLU) | (None, 32, 224, 224) | 0 | convolution2d_2[0][0] |
| ----- | | | |
| maxpooling2d_2 (MaxPooling2D) | (None, 32, 112, 112) | 0 | leakyrelu_2[0][0] |
| ----- | | | |
| convolution2d_3 (Convolution2D) | (None, 64, 112, 112) | 18496 | maxpooling2d_2[0][0] |
| ----- | | | |
| leakyrelu_3 (LeakyReLU) | (None, 64, 112, 112) | 0 | convolution2d_3[0][0] |

| | | | |
|---------------------------------|---------------------|---------|-----------------------|
| maxpooling2d_3 (MaxPooling2D) | (None, 64, 56, 56) | 0 | leakyrelu_3[0][0] |
| convolution2d_4 (Convolution2D) | (None, 128, 56, 56) | 73856 | maxpooling2d_3[0][0] |
| leakyrelu_4 (LeakyReLU) | (None, 128, 56, 56) | 0 | convolution2d_4[0][0] |
| maxpooling2d_4 (MaxPooling2D) | (None, 128, 28, 28) | 0 | leakyrelu_4[0][0] |
| convolution2d_5 (Convolution2D) | (None, 256, 28, 28) | 295168 | maxpooling2d_4[0][0] |
| leakyrelu_5 (LeakyReLU) | (None, 256, 28, 28) | 0 | convolution2d_5[0][0] |
| maxpooling2d_5 (MaxPooling2D) | (None, 256, 14, 14) | 0 | leakyrelu_5[0][0] |
| convolution2d_6 (Convolution2D) | (None, 512, 14, 14) | 1180160 | maxpooling2d_5[0][0] |
| leakyrelu_6 (LeakyReLU) | (None, 512, 14, 14) | 0 | convolution2d_6[0][0] |
| maxpooling2d_6 (MaxPooling2D) | (None, 512, 7, 7) | 0 | leakyrelu_6[0][0] |
| convolution2d_7 (Convolution2D) | (None, 1024, 7, 7) | 4719616 | maxpooling2d_6[0][0] |
| leakyrelu_7 (LeakyReLU) | (None, 1024, 7, 7) | 0 | convolution2d_7[0][0] |

| | | | |
|---------------------------------|--------------------|----------|-----------------------|
| convolution2d_8 (Convolution2D) | (None, 1024, 7, 7) | 9438208 | leakyrelu_7[0][0] |
| <hr/> | | | |
| leakyrelu_8 (LeakyReLU) | (None, 1024, 7, 7) | 0 | convolution2d_8[0][0] |
| <hr/> | | | |
| convolution2d_9 (Convolution2D) | (None, 1024, 7, 7) | 9438208 | leakyrelu_8[0][0] |
| <hr/> | | | |
| leakyrelu_9 (LeakyReLU) | (None, 1024, 7, 7) | 0 | convolution2d_9[0][0] |
| <hr/> | | | |
| flatten_1 (Flatten) | (None, 50176) | 0 | leakyrelu_9[0][0] |
| <hr/> | | | |
| dense_1 (Dense) | (None, 256) | 12845312 | flatten_1[0][0] |
| <hr/> | | | |
| dense_2 (Dense) | (None, 4096) | 1052672 | dense_1[0][0] |
| <hr/> | | | |
| leakyrelu_10 (LeakyReLU) | (None, 4096) | 0 | dense_2[0][0] |
| <hr/> | | | |
| dense_3 (Dense) | (None, 1470) | 6022590 | leakyrelu_10[0][0] |
| <hr/> | | | |
| ===== | | | |
| ===== | | | |

Total params: 45,089,374

Trainable params: 45,089,374

Non-trainable params: 0

Load Pre-trained Yolo weights and Train one image

Load pretrained weights of yolo provided from

https://drive.google.com/file/d/0B1tW_VtY7onibmdQWE1zVERxcjQ/view


Data output for bounding box and probability of each class which recorded in "data.md" file.

The bounding box of vehicle in training image:

```
[-6.2443885803222656, 22.23392595563616, 2177.4734942449722, 1823.8973605994979]  
[8.1763703482491632, 2.4748785836356029, 420.87584318220615, 633.60562289156951]  
[-7.6249623979840964, -4.1853103637695312, 3655.9590410698584, 11427.779967978189]  
[39.074292864118306, 5.6808046613420755, 42097.977821622044, 32.97495707947769]  
[-3.2266101837158203, -8.3505581447056354, 3163.0630805657856, 3610.5449726496736]  
[-13.818937029157366, 18.408437456403458, 213.93004241378458, 5060.5619251030148]  
[1.5318431173052107, -5.3705831255231589, 371.11749337709989, 10926.247901296301]  
[-0.36035473006112234, -3.1194964817592075, 137.45633478127365, 16792.321420513326]  
[-18.485462733677455, 24.96602303641183, 31048.148896806175, 263.02700191329131]  
[18.209995814732142, -6.85974611554827, 4270.4358333386481, 650.95543337824347]  
[4.1464734758649557, -19.608760288783483, 403.16572382865706, 2888.5655044692103]  
[10.283409118652344, 11.294999258858818, 197.08981332987605, 22.75501516451186]  
[10.924356733049665, 4.104710715157645, 1433.11478402793, 4985.7506377550308]  
[-2.2601637159075056, 6.2452278137207031, 18.584896667613293, 2015.9841025611968]  
[14.397278921944755, 6.0225312369210382, 12340.756753989495, 2348.5042427848239]  
[21.816515241350448, -6.2865398951939175, 5.1132214578594812, 28.415502496536874]  
[-6.5367927551269531, 10.683708190917969, 18.260654216897819, 9512.4072314335499]  
[3.9819752829415456, 0.15059631211417063, 37338.503499800572, 2321.6925270596985]  
[-35.955019269670757, 5.9842344011579245, 2445.5093257385452, 14.417907359584433]
```

Draw bounding box in image

Result:



Pipeline for video stream

Generate "project_output.mp4"